

Lecture 23: Application Layer

Design Patterns, DNS, HTTP, and QUIC

EC 441 — Introduction to Computer Networking
Boston University

Spring 2026

Session 1 — Patterns and DNS

- Design patterns of app protocols
- Client-server vs P2P (BitTorrent)
- DNS: the biggest database on earth
- `dig` + `trace` live

Session 2 — HTTP and Friends

- HTTP/1.1 — the debuggable baseline
- HTTP/2 — binary, multiplexed
- HTTP/3 / QUIC — over UDP
- Email + webmail complication
- WebSocket callout

Theme

Last week you built from sockets. Today: the protocols people actually write at that layer.

How many application-layer protocols is your laptop speaking *right now*?

- DNS (every name you resolve)
- HTTPS (web)
- QUIC (HTTP/3 on Google, Cloudflare, YouTube)
- Mail fetch (IMAP/POP) if a desktop client is open
- Push notifications (APNS / FCM, over TLS)
- SSH if a terminal is open
- NTP (your clock)

The application layer is where protocol design has the most freedom — and where almost all of the protocols you directly use live.

Axes of Application Protocol Design

Axis	Option A	Option B
Interaction	request/response	push / pub-sub
State	stateless	stateful
Encoding	text (ASCII)	binary
Direction	pull	push
Connection	short-lived	persistent
Architecture	client-server	peer-to-peer

Key skill: when you meet a new protocol, name it along these axes. You'll understand it faster.

Examples by Pattern

Protocol	Interaction	State	Encoding	Connection
HTTP/1.1	request/resp	stateless	text	semi-persistent
HTTP/2	request/resp	stateless	binary	persistent
DNS	request/resp	stateless	binary	short-lived (UDP)
SSH	stream	stateful	binary	long-lived
SMTP	command/resp	session	text	short-lived
WebSocket	duplex	stateful	framed	long-lived
BitTorrent	P2P	swarm	binary	many short

Stateless: The HTTP Cookie

HTTP is *stateless*: each request stands alone, the server does not remember you between requests.

How does “you are logged in” work, then?

The client carries the state and presents it on every request.

- Server sets Set-Cookie: `session=abc123` on login response.
- Client stores it and sends Cookie: `session=abc123` on every subsequent request.
- Server looks up abc123 in its session table.

Stateless on the wire; state lives in the cookie jar and the server’s database. This shift (state → client) is the single most important pattern in web architecture.

Two Families

Client-server — asymmetric roles

- Server listens, waits. Client initiates.
- HTTP, DNS, SSH, SMTP, IMAP.
- Scale: replication (CDNs), anycast, load balancers.

Peer-to-peer — symmetric roles

- Every node is both client and server.
- No central coordinator, or only a lightweight one.
- Canonical example: **BitTorrent**.

BitTorrent: Mechanism Design in a Protocol

Problem: distribute a big file to many recipients without paying for all the bandwidth from one server.

Design beats:

- ① **Chunks with hashes.** File split into fixed-size chunks; each chunk has a cryptographic hash in the .torrent metadata. You can verify any chunk from any peer — no need to trust peers.
- ② **Peer discovery.** Trackers (hybrid) or DHT (fully decentralized).
- ③ **Rarest-first chunk selection.** Keep the rare chunks alive in the swarm.
- ④ **Tit-for-tat choking.** Prefer peers who upload back to you; throttle free-riders.

Popularity increases capacity. That is the P2P win.

Why P2P Doesn't Win Everywhere

P2P is a great fit for:

- Bulk file distribution (Linux ISOs, game patches, datasets).
- Censorship-resistant storage (IPFS).
- Consensus protocols (blockchain).

P2P is a bad fit for:

- Anything needing a single source of truth (account balance, authoritative DNS, email delivery).
- Real-time interactive services (chat, video call).
- Billing (who is the customer?).

Contrast with DNS. DNS is a *tree with delegation*, not a mesh. Hierarchical, not P2P. Yet it is one of the most decentralized real systems we have.

DNS: The Namespace

Names are a tree, read right to left:

```
www.eng.bu.edu.  
|   |   |   |  
|   |   |   +-- TLD: edu  
|   |   +----- second level: bu  
|   +----- subdomain: eng  
+----- host: www
```

The trailing . is the root.

- **Root:** 13 server *identities*, hundreds of physical servers via anycast.
- **TLDs:** edu, com, org, country-code (.uk, .jp), new gTLDs.
- **Authoritative:** operated by the domain owner (or their DNS provider: Cloudflare, AWS Route 53).

DNS Resolution Chain

Your laptop wants `www.eng.bu.edu`:

- ① Stub resolver on laptop → **recursive resolver** (1.1.1.1 / 8.8.8.8 / your ISP).
- ② Resolver cache hit? Done.
- ③ Otherwise walk:
 - root → “who handles edu?”
 - edu → “who handles bu.edu?”
 - bu.edu → “who handles eng.bu.edu?”
 - eng.bu.edu → “what is www.eng.bu.edu?”
- ④ Answer bubbles back to the laptop.

Recursive at the resolver. Iterative at authoritative servers. Most laptops never speak directly to a root or TLD server.

DNS Record Types

Type	Purpose
A	IPv4 address
AAAA	IPv6 address (four A's — four times as long)
CNAME	Alias to another name
MX	Mail exchange, with priority
NS	Name servers for a zone
TXT	Free-form text; SPF, DKIM, domain verification
SOA	Start of authority; zone metadata, TTLs

Caching and TTLs make DNS scale.

- Web records: TTL $\sim$ 300–3600 s.
- Root hints: TTL $\sim$ 6 days.
- DNS handles $\sim 10^{12}$ queries/day without answering them all fresh.

```
dig www.eng.bu.edu +short      # just the IP
dig www.eng.bu.edu             # full output
dig bu.edu MX +short           # mail exchanges
dig bu.edu NS +short           # name servers
dig +trace www.eng.bu.edu       # walk root -> TLD -> auth
dig -x 8.8.8.8 +short          # reverse lookup (PTR)
```

See: `demo_dig_trace_l23.sh`

In +trace output: watch for `root-servers.net` at the top, then the `edu` servers, then `bu.edu`, then the final `A` record.

Transport	Port	Use
UDP	53	Default. Small queries, small replies.
TCP	53	Large responses, zone transfers, DNSSEC.
TLS	853	DoT — DNS over TLS. Privacy.
HTTPS	443	DoH — DNS over HTTPS. Privacy, blends with web traffic.

Why DoH is controversial

DoH looks like HTTPS on port 443, so network operators cannot distinguish it from web traffic. Good for privacy, hard for local policy (enterprise filtering, parental controls, captive portals).

What Is “The Web”?

HTTP fetches **one** resource. One URL, one response. That's it.

“The Web” is what emerges when HTML documents *link to other resources*:

- `<a href=“...”>` — follow a link
- `<img src=“...”>` — embed an image
- `<link rel=“stylesheet” href=“...”>` — pull in CSS
- `<script src=“...”>` — pull in JavaScript
- `<video>`, `<iframe>`, `<object>`, fonts, tracking pixels, ...

The browser is the coordinator. Parse HTML → discover links → fire off more HTTP requests → render.

A modern page: **1 HTML document + hundreds of sub-resources.** This is the workload HTTP/2 and HTTP/3 are built for.

MIME Types: What Are These Bytes?

When the browser gets a response, how does it know *what the bytes mean*?

The Content-Type header carries a **MIME type**:

MIME type	Used for
text/html	HTML documents
text/css	Stylesheets
application/javascript	JavaScript
application/json	JSON (APIs)
image/png, image/jpeg, image/webp	Images
application/pdf	PDF documents
video/mp4	Video
application/octet-stream	Unknown binary (“download me”)

Originally designed for email (**M**ultipurpose **I**nternet **M**ail **E**xtensions); now universal.

Live: `curl -sI https://example.com | grep -i content-type`

HTTP/1.1 — The Baseline

```
GET /index.html HTTP/1.1
Host: example.com
User-Agent: curl/8.4.0
Accept: */*
```

Response:

```
HTTP/1.1 200 OK
Content-Type: text/html; charset=UTF-8
Content-Length: 1256
Connection: keep-alive

<!DOCTYPE html>...
```

ASCII the whole way. You can telnet `example.com 80` and type it by hand.

HTTP/1.1 — Key Features

- **Methods:** GET, POST, PUT, DELETE, HEAD, OPTIONS, PATCH.
- **Status codes:** 1xx info, 2xx ok, 3xx redirect, 4xx client error, 5xx server error.
- **Headers:** key-value, extensible, case-insensitive.
- **Body framing:** Content-Length or Transfer-Encoding: chunked.
- **Stateless** — cookies provide the session illusion.
- **Keep-alive:** reuse one TCP connection across requests.

HOL problem. Pipelining (multiple requests on one connection) required responses to come back in request order. One slow response $\Rightarrow$ everything behind it stalls. Browsers gave up and opened 6+ parallel connections.

Live: curl with HTTP/1.1

```
curl -v --http1.1 https://example.com

> GET / HTTP/1.1
> Host: example.com
> User-Agent: curl/8.4.0
>
< HTTP/1.1 200 OK
< Content-Type: text/html
< Content-Length: 1256
...
```

See: `demo_curl_versions_l23.sh`

Lines starting with > are what curl sends; < is what the server returns; * is curl's connection state.

HTTP/2 — Binary and Multiplexed

Same semantics (methods, headers, status codes). New wire format.

- **Binary framing.** Every message is a stream of binary frames. No more telnet debugging.
- **Multiplexing on one TCP connection.** Many streams in flight at once; frames interleave. **App-level HOL blocking solved.**
- **HPACK header compression.** Don't re-send the same headers on every request.
- **Server push** — proactively send resources. Mostly deprecated in practice.

What HTTP/2 does *not* solve

TCP-layer HOL blocking. A single lost TCP segment stalls every stream, because TCP is a byte stream that must be delivered in order.

Live: curl with HTTP/2

```
curl -v --http2 -o /dev/null https://www.cloudflare.com

* using HTTP/2
* h2h3 [:method: GET]
* h2h3 [:path: /]
* h2h3 [:authority: www.cloudflare.com]
...
< HTTP/2 200
< content-type: text/html
...
```

:method, :path, :authority are HTTP/2 *pseudo-headers* — what Host, the method, and the path became in the binary wire format.

HTTP/3 runs over QUIC. QUIC runs over UDP.

Why UDP?

- TCP cannot be meaningfully changed in every middlebox on the internet.
- NATs, firewalls, load balancers make strong assumptions about TCP.
- UDP is dumb enough to mostly be left alone.
- Move the reliable transport logic into **user space**, built on UDP.

What QUIC provides

- Reliable, ordered delivery **per stream** (TCP HOL *gone*).
- Multiplexed independent streams.
- Built-in TLS 1.3 in the handshake (not a layer on top).
- 0-RTT resumption on reconnect.
- Connection migration across IP changes.

QUIC: Callbacks to Earlier Lectures

- **L18 introduced QUIC** as the reason UDP gets a second act. *This is it.*
- **L20 congestion control.** QUIC runs its own CUBIC / BBR in user space. Same algorithms, different implementation surface.
- **L22 sockets.** QUIC rides on a UDP socket; everything else is library code.

The big picture

QUIC took what kernels do and put it in a library. A browser update ships a new transport to a billion users overnight — no kernel patch needed.

Live: curl with HTTP/3

```
curl -v --http3 -o /dev/null https://www.cloudflare.com  
  
* Connected to www.cloudflare.com:443 via HTTP/3  
* using HTTP/3  
...  
< HTTP/3 200
```

Requires a curl with HTTP/3 support. Check with `curl --version` — look for HTTP3 in the Features line.

On macOS: `brew install curl` gives a modern build. The system curl may not include HTTP/3.

HTTP Timeline

Year	Version	Transport	Wire	Fixes
1996	HTTP/1.0	TCP	text	—
1999	HTTP/1.1	TCP	text	keep-alive, host header
2015	HTTP/2	TCP	binary	multiplexing, HPACK
2022	HTTP/3	QUIC (UDP)	binary	TCP HOL gone, TLS built-in

Each version keeps HTTP *semantics* (methods, status, headers). The wire format, transport, and connection model keep changing.

Email Is Three Protocols

Protocol	Port(s)	Role
SMTP	25 / 587 / 465	send (client → server, server → server)
IMAP	143 / 993	read mailboxes on a server
POP3	110 / 995	download mailboxes (legacy)

Classical picture: desktop mail client speaks SMTP to send, IMAP to read.

Classical design patterns: text-based, command/response, stateful-within-session, short-lived connections. Readable with `telnet` — even easier to debug than HTTP/1.1.

The Webmail Complication

Gmail is not speaking SMTP or IMAP to your browser.

- Your browser speaks **HTTPS** to a Google web app.
- Google's servers speak SMTP and IMAP *on your behalf* internally — to Google's own mail storage, and to external mail servers (outlook.com, bu.edu) for delivery.
- The user-facing protocol is JSON-over-HTTPS, not email.

Consequence: most users never see SMTP or IMAP.

- Every webmail provider ships its own UI over its own private HTTP API.
- The mail protocols between servers still run — and still carry massive backward-compatibility baggage.
- This is one big reason email evolves slowly.

WebSocket — The Escape Hatch

Some apps don't fit request/response: chat, live dashboards, collaborative editing, multiplayer games.

WebSocket (RFC 6455): start with an HTTP *upgrade*, then switch to a persistent duplex frame channel.

```
GET /chat HTTP/1.1
Host: example.com
Upgrade: websocket
Connection: Upgrade
Sec-WebSocket-Key: ...
```

After the 101 Switching Protocols response, the connection stops being HTTP. Both ends send messages at any time.

Used by: Slack, Google Docs, Discord, live stock feeds.

The HTML5 connection. WebSocket is defined in the **HTML5 Living Standard** (WHATWG). Browsers expose it as `new WebSocket(url)` in JavaScript — the browser API lives in HTML5; the wire protocol is RFC 6455.

Full Circle: Back to Lecture 2

Three questions worth asking about every modern web page:

- Why is "Hello <two Chinese chars>" in HTML **12 bytes**, not 24?
- Why is a 2-megapixel photo **100 KB** JPEG, not 6 MB?
- Why is 1080p Netflix **5 Mb/s**, not 1.5 Gb/s?

Answer: Shannon (L2). Every MIME type carries an encoding whose job is to push the bit count toward the entropy.

Content	MIME type	L2 said
HTML text	text/html; charset=UTF-8	UTF-8: ASCII compat, 1–4 B/char
Stylesheet	text/css	gzip/Brotli in transit
JavaScript	application/javascript	gzip/Brotli in transit
Photograph	image/jpeg	6.2 MB → 100 KB (60×)
Video stream	video/mp4	1.49 Gb/s → 5 Mb/s (300×)

Without L2, none of this works. Uncompressed video alone would saturate a 1 Gb/s Ethernet link for one stream.

One HTTP Response, Every Layer

What happens when you load *one* frame of a Netflix stream:

Layer	What it contributes
L2 — information	H.264 encodes frame: $\sim 1.49 \text{ Gb/s} \rightarrow 5 \text{ Mb/s}$
L3 — physical	Bits on fiber / radio / copper
L4 — wireless	OFDM subcarriers on 802.11ax
L5–L6 — error coding	CRC, FEC protect each frame
L7 — MAC	CSMA/CA arbitrates the air
L8 — Ethernet	Frame with preamble, MAC, FCS
L9 — WiFi	802.11 frame, probe/auth/assoc done
L13–L17 — network	IP datagram, TTL, routing, NAT
L18–L20 — transport	TCP (or QUIC) segment, cwnd, RTT
L21–L22 — tools	ss, tcpdump, sockets, Scapy, Mininet
L23 — application	HTTP/2 stream carrying video/mp4
→ screen	One decoded 1080p frame. 33 ms later, the next.

The whole semester lives in one HTTP response.

What tcpdump Could Read Today

Every protocol in this lecture runs in **cleartext** on the wire by default:

- HTTP/1.1 headers, cookies, bodies — all visible.
- DNS queries — visible to your ISP, your WiFi AP owner, governments.
- SMTP between servers — often encrypted via STARTTLS, but originally plaintext.

QUIC and DoH push encryption into defaults. The big answer is **HTTPS** = HTTP + **TLS**.

L24: Cryptography and Security

What is inside TLS? Symmetric crypto, asymmetric crypto, hashes, digital signatures, certificates, and a global PKI. That is L24.

What's Next

Lecture	Date	Topic
L24	Tue Apr 28	Cryptography and security — $\alpha/\beta/\tau$ (actors), TLS
L25	Thu Apr 30	Project demo day

Closing line

Everything you watched today, tcpdump could read. Next lecture: how we fix that.

Acronyms

Acronym	Expansion
CDN	Content Delivery Network
DNS	Domain Name System
DoH	DNS over HTTPS
DoT	DNS over TLS
HTML	HyperText Markup Language
HTTP	HyperText Transfer Protocol
HTTPS	HTTP Secure
IMAP	Internet Message Access Protocol
ISP	Internet Service Provider
MIME	Multipurpose Internet Mail Extensions

Acronym	Expansion
P2P	Peer-to-Peer
POP	Post Office Protocol
QUIC	Quick UDP Internet Connections
SMTP	Simple Mail Transfer Protocol
SSL	Secure Sockets Layer
TCP	Transmission Control Protocol
TLD	Top-Level Domain
TLS	Transport Layer Security
TTL	Time To Live
UDP	User Datagram Protocol
URL	Uniform Resource Locator